

# logistic-candy

March 21, 2024

```
[1]: import pandas as pd
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: df = pd.read_csv("logistic-candy-data.csv")
df
```

```
[2]:
```

	competitorname	chocolate	fruity	caramel	peanutyalmondy	\
0	100 Grand	1	0	1	0	
1	3 Musketeers	1	0	0	0	
2	One dime	0	0	0	0	
3	One quarter	0	0	0	0	
4	Air Heads	0	1	0	0	
..	...	...	...	...	...	
80	Twizzlers	0	1	0	0	
81	Warheads	0	1	0	0	
82	Welch's Fruit Snacks	0	1	0	0	
83	Werther's Original Caramel	0	0	1	0	
84	Whoppers	1	0	0	0	

  

	nougat	crispedricewafer	hard	bar	pluribus	sugarpercent	pricepercent	\
0	0	1	0	1	0	0.732	0.860	
1	1	0	0	1	0	0.604	0.511	
2	0	0	0	0	0	0.011	0.116	
3	0	0	0	0	0	0.011	0.511	
4	0	0	0	0	0	0.906	0.511	
..	...	...	...	...	...	...	...	
80	0	0	0	0	0	0.220	0.116	
81	0	0	1	0	0	0.093	0.116	
82	0	0	0	0	1	0.313	0.313	
83	0	0	1	0	0	0.186	0.267	
84	0	1	0	0	1	0.872	0.848	

  

	winpercent
0	66.971725
1	67.602936
2	32.261086
3	46.116505

```

4      52.341465
..      ...
80     45.466282
81     39.011898
82     44.375519
83     41.904308
84     49.524113

```

[85 rows x 13 columns]

```
[3]: df.shape
```

```
[3]: (85, 13)
```

```
[4]: df.isnull().sum()
```

```

[4]: competitorname    0
     chocolate         0
     fruity            0
     caramel           0
     peanutyalmondy    0
     nougat            0
     crispedricewafer  0
     hard              0
     bar               0
     pluribus          0
     sugarpercent      0
     pricepercent      0
     winpercent        0
     dtype: int64

```

```
[5]: df['chocolate'].value_counts()
```

```

[5]: chocolate
0      48
1      37
     Name: count, dtype: int64

```

```

[6]: x = df.drop(columns=["competitorname", "chocolate", "fruity", "caramel",
    ↪ "peanutyalmondy", "nougat", "crispedricewafer", "hard", "bar", "pluribus"],
    ↪ axis=1)
     x

```

```

[6]:      sugarpercent  pricepercent  winpercent
0           0.732         0.860    66.971725
1           0.604         0.511    67.602936
2           0.011         0.116    32.261086
3           0.011         0.511    46.116505

```

4	0.906	0.511	52.341465
..	...	...	...
80	0.220	0.116	45.466282
81	0.093	0.116	39.011898
82	0.313	0.313	44.375519
83	0.186	0.267	41.904308
84	0.872	0.848	49.524113

[85 rows x 3 columns]

```
[7]: y = df["chocolate"]
      y
```

```
[7]: 0    1
      1    1
      2    0
      3    0
      4    0
      ..
      80   0
      81   0
      82   0
      83   0
      84   1
      Name: chocolate, Length: 85, dtype: int64
```

```
[8]: from sklearn.model_selection import train_test_split
```

```
[9]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
      ↪random_state=42, stratify=y)
```

```
[10]: from sklearn.linear_model import LogisticRegression
```

```
[11]: model = LogisticRegression(max_iter=1000).fit(x_train, y_train)
      model
```

```
[11]: LogisticRegression(max_iter=1000)
```

```
[12]: y_pred = model.predict(x_test)
      y_pred
```

```
[12]: array([0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1], dtype=int64)
```

```
[13]: from sklearn.metrics import accuracy_score
```

```
[14]: accuracy = accuracy_score(y_test, y_pred)
```

```
[15]: print(f"Accuracy = ", accuracy)
```

```
Accuracy = 0.8823529411764706
```

## MODEL OPTIMIZATION

```
[16]: from sklearn.model_selection import GridSearchCV
```

```
[17]: model = LogisticRegression()  
model
```

```
[17]: LogisticRegression()
```

```
[18]: param_grid = {  
    'penalty': ['l2', None],  
    'solver': ['liblinear', 'newton-cg', 'newton-cholesky', 'sag', 'saga'],  
    'C': [1.0, 1.5]  
}
```

```
[19]: grid_search = GridSearchCV(model, param_grid, cv=5, n_jobs=-1)  
grid_search.fit(x_train, y_train)
```

```
[19]: GridSearchCV(cv=5, estimator=LogisticRegression(), n_jobs=-1,  
    param_grid={'C': [1.0, 1.5], 'penalty': ['l2', None],  
    'solver': ['liblinear', 'newton-cg', 'newton-cholesky',  
    'sag', 'saga']})
```

```
[20]: best_params = grid_search.best_params_  
print("Best Parameters :", best_params)
```

```
Best Parameters : {'C': 1.0, 'penalty': 'l2', 'solver': 'liblinear'}
```

```
[21]: best_model = LogisticRegression(**best_params)  
best_model.fit(x_train, y_train)  
best_model
```

```
[21]: LogisticRegression(solver='liblinear')
```

```
[22]: y_pred = best_model.predict(x_test)  
y_pred
```

```
[22]: array([0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1], dtype=int64)
```

```
[23]: accuracy = accuracy_score(y_test, y_pred)
```

```
[24]: print("Best Parameters :", best_params)  
print(f"Accuracy = ", accuracy)
```

```
Best Parameters : {'C': 1.0, 'penalty': 'l2', 'solver': 'liblinear'}  
Accuracy = 0.8823529411764706
```

```
[ ]:
```